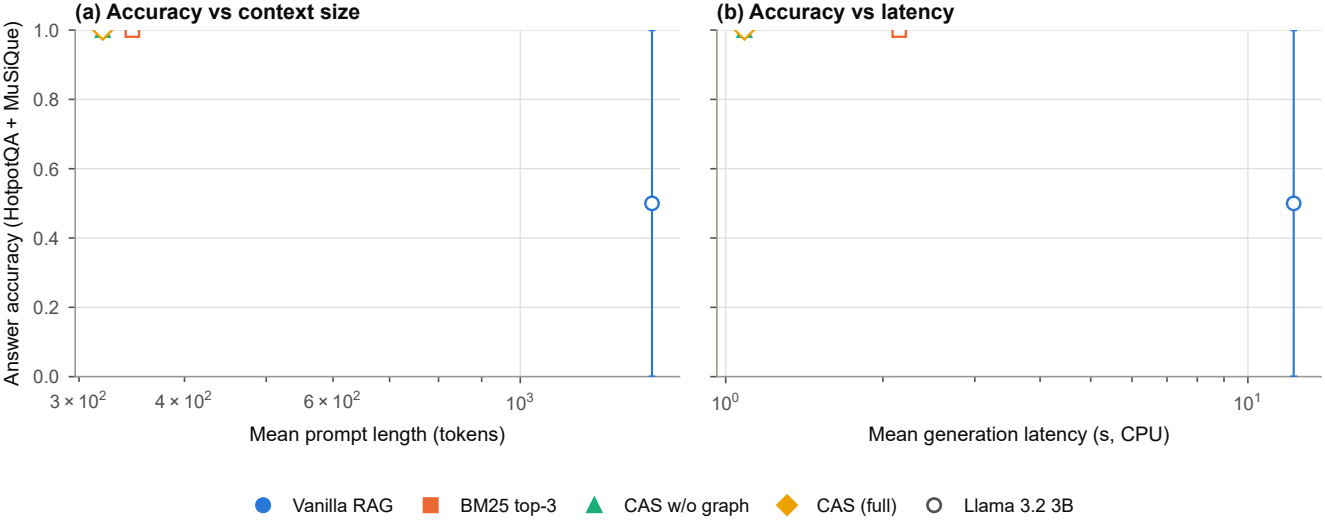
(a) Accuracy vs context size
(b) Accuracy vs latency
Answer accuracy (HotpotQA + MuSiQue)
1.0
0.8
0.6
0.4
0.2
0.0
3 × 10^2
4 × 10^2
6 × 10^2
10^3
Mean prompt length (tokens)
10^0
10^1
Mean generation latency (s, CPU)
Vanilla RAG
BM25 top-3
CAS w/o graph
CAS (full)
Llama 3.2 3B